KukuFM PromoFinder solution

I rank screenplay moments for editor review.
Problem: I rank 20 to 60 second screenplay moments with an early hook and unresolved ending.

Assumptions: screenplay timestamps align with video. An editor checks every boundary. AI judges compact excerpts. Code parses files and validates quotes, beat edges, and duration.
Run: 50 episodes, 30 promos, 39 rejected.
Tokens: initial live pass 173814, current run 0, 173814 measured total.
Average per episode: 3476.28 tokens. Ten-show projection: 1738140 tokens.
Configured cost estimate: $0.0000 for this run, $0.0000 for ten similar shows. Price settings are not provider billing records.
Runtime: 15.4 seconds.
Quality results:
validity: 30/30 (1.000)
quote_verification: 30/30 (1.000)
duration_compliance: 30/30 (1.000)
snap_rate: 30/30 (1.000)
raw_proposal_validity: 12/300 (0.040)
rejected_share: 39/69 (0.565)
cold_viewer_conflict_missing: 0/0

Blind labels: 0.
Failures and tools:
Rule-valid output does not guarantee viewer appeal. A weak candidate can pass every mechanical check while its opening relies on visual action that has no strong first-line hook, or while its ending is merely dramatic rather than unresolved.
The largest known failure was a full live ranking error at global ranking, recorded as an `UnboundLocalError` after the first live pass. The fallback pass and export path have since been tested offline.
The 30 output promos therefore include candidates that need human review. No blind-label study was completed in the submission window. The current quality threshold is conservative: every exported promo must pass duration, evidence, boundary, and overlap checks, but the editor decides cultural fit and social appeal.
Tools and one AI mistake:
I used Python, Streamlit, FastAPI, ReportLab, Pytest, Groq, GitHub, Vercel, and Streamlit Community Cloud. Groq was used for compact episode judgments and stability checks; deterministic Python handled parsing and validation. All AI output is judged against timestamps and screenplay text rather than accepted by display alone.
One observed model mistake was a Groq proposal that used valid JSON and accurate timestamps but changed the exact screenplay wording in its quote evidence. The quote verifier rejected that proposal. This showed why exact quote verification must stay in code outside the model.